

# Table of Contents

- [1 Introduction](#)
- [2 Data](#)
  - [2.1 Load the data](#)
  - [2.2 Sampling: Train Test Split](#)
    - [2.2.1 Random Sampling](#)
    - [2.2.2 Stratified Sampling](#)
    - [2.2.3 Random vs Stratified](#)
- [3 Exploratory Data Analysis](#)
  - [3.0.1 Data Types](#)
    - [3.0.1.1 Convert datetime column](#)
  - [3.0.2 Value counts in each column \(categorical\)](#)
- [4 Data Transformation](#)
  - [4.1 Missing Values](#)
    - [4.1.1 Precipitation column](#)
    - [4.1.2 Road Surface Conditions column](#)
    - [4.1.3 Transform the test dataset](#)
  - [4.2 Feature Engineering](#)
  - [4.3 Create dummy variables](#)
  - [4.4 Finalise trainset for analysis](#)
- [5 Visualisations](#)
  - [5.0.1 Heatmap](#)
  - [5.0.2 Average Hourly 2018](#)
  - [5.0.3 Plotting a histogram](#)
  - [5.0.4 Average Monthly 2018](#)
  - [5.0.5 Time Series](#)
- [6 Correlations](#)
- [7 Time Series Model](#)
  - [7.0.1 Additional Preprocessing and Transformation](#)
    - [7.0.1.1 Decompose time series](#)
    - [7.0.1.2 Testing for stationarity](#)
    - [7.0.1.3 Log Transformation](#)
    - [7.0.1.4 Differencing](#)
  - [7.1 Baseline](#)
  - [7.2 ARIMA](#)
    - [7.2.1 Determine the parameters](#)
  - [7.3 Hyperparameter Tuning](#)
    - [7.3.1 Build ARIMA Model](#)
    - [7.3.2 Evaluate ARIMA on test data](#)
- [8 Further Models](#)
  - [8.1 ARIMAX](#)
    - [8.1.1 Add exogenous variables](#)
    - [8.1.2 Build ARIMAX model](#)
      - [8.1.2.1 Evaluate model on test data](#)
- [9 Model Evaluation](#)
  - [9.1 RMSE Results](#)
  - [9.2 MDA Results](#)
- [10 Conclusion](#)

# Introduction

Road assistance companies need to efficiently plan ahead of time to allocate their resources to meet demand on an hourly basis. Without meeting the demand, this could cause unnecessary prolonged obstructions on roads, leading to traffic and in effect resulting increased fuel consumption and effect on quality of life and economic productivity.

This project will involve analysing the hourly number of accidents over the course of 2018 to find trends in number of accidents around a specific hour of the day; also check factors like urban or rural area and precipitation affect the number of accidents that occur.

Import the libraries required for analysis

In [1]:

```
import warnings
warnings.filterwarnings("ignore")

import time
#library for data manipulation and preprocessing
import pandas as pd

#library for number crunching
import numpy as np

#library for visualisation
import matplotlib.pyplot as plt
import statsmodels.api as sm
import scipy.stats
import seaborn as sns
import matplotlib.dates as mdates
plt.style.use('bmh')

from sklearn.metrics import mean_squared_error
#time series data
from sklearn.model_selection import TimeSeriesSplit, GridSearchCV

def get_mda(y, yhat):
    """Mean Directional Accuracy, as per:
    https://www.wikiwand.com/en/Mean\_Directional\_Accuracy
    """
    a = np.sign(np.diff(y))
    b = np.sign(np.diff(yhat))
    return np.sum(a == b)/a.shape[0]
```

## Data

### Load the data

In [2]:

```
df = pd.read_csv("UK Road Accidents 2018.csv")
df.head(5)
```

Out[2]:

	datetime	Count	Day_of_Week	Precipitation	High_Winds	Road_Surface_Conditions	Light_
0	2018-01-01 00:00:00	18	Monday	Fine	False	Dry	Darkr
1	2018-01-01 01:00:00	9	Monday	Fine	False	Wet or damp	Darkr
2	2018-01-01 02:00:00	14	Monday	Fine	False	Wet or damp	Darkr
3	2018-01-01 03:00:00	10	Monday	Fine	False	Wet or damp	Darkr
4	2018-01-01 04:00:00	8	Monday	Fine	False	Dry	Darkr

In [3]:

```
df.shape
```

Out[3]:

```
(8540, 8)
```

There are 8 columns in total in the dataset with 8540 entries. Each row contains records on accident occurrences in a year, month, day, time, day of the week, weather percipitations, high winds, road surface conditions, light conditions and whether accident occurred in an urban or rural area.

The count column is the target variable in the dataset that will be predicted in the predictive models.

## Sampling: Train Test Split

A train-test split will be performed using both `random sampling` and `stratified sampling` methods; where most appropriate will be chosen to ensure that the critical attributes of a population are correctly represented.

## Random Sampling

In [4]:

```
from sklearn.model_selection import train_test_split

train_set, test_set = train_test_split(df, test_size=0.2, random_state=7, shuffle=False)
print(f"{train_set.shape[0]} train instances and {test_set.shape[0]} test instances")
```

6832 train instances and 1708 test instances

## Stratified Sampling

In [5]:

```
from sklearn.model_selection import StratifiedShuffleSplit

stratified_splitter = StratifiedShuffleSplit(n_splits=1, test_size=0.2, random_state=7)

train_index, test_index = list(stratified_splitter.split(df, df["Day_of_Week"]))
strat_train_set = df.loc[train_index]
strat_test_set = df.loc[test_index]
```

## Random vs Stratified

In [6]:

```
def day_of_week_proportions(data):
    return data["Day_of_Week"].value_counts() / len(data)

# create a random split
rand_train_set, rand_test_set = train_test_split(df, test_size=0.2, random_state=7)

# create a temporary dataframe for easy visualization
df_tmp = pd.DataFrame({
    "Overall": day_of_week_proportions(df),
    "Random test set": day_of_week_proportions(rand_test_set),
    "Stratified test set": day_of_week_proportions(strat_test_set),
}).sort_index()

# add two columns for the percent of the difference to the overall proportion
df_tmp["Rand. %error"] = 100 * df_tmp["Random test set"] / df_tmp["Overall"] - 100
df_tmp["Strat. %error"] = 100 * df_tmp["Stratified test set"] / df_tmp["Overall"] - 100

df_tmp
```

Out[6]:

	Overall	Random test set	Stratified test set	Rand. %error	Strat. %error
<b>Friday</b>	0.142506	0.151639	0.142272	6.409203	-0.164339
<b>Monday</b>	0.144496	0.149297	0.144614	3.322528	0.081037
<b>Saturday</b>	0.145550	0.144614	0.145785	-0.643604	0.160901
<b>Sunday</b>	0.145667	0.144614	0.145785	-0.723473	0.080386
<b>Thursday</b>	0.140984	0.140515	0.141101	-0.332226	0.083056
<b>Tuesday</b>	0.139578	0.131148	0.139344	-6.040268	-0.167785
<b>Wednesday</b>	0.141218	0.138173	0.141101	-2.155887	-0.082919

As we can see in the above table, the random splitting produces a test set where Friday and Monday are over represented by 6% and 3% respectively. Tuesday and Wednesday are under-represented by 6% and 2% respectively.

Stratified sampling resulted in under- or over-representations of the days of the week by no more than 0.16%. Choose to use stratified sampling as it is more representative of the dataset.

Rename variables

In [7]:

```
trainset = strat_train_set
testset = strat_test_set
```

# Exploratory Data Analysis

## Data Types

In [8]:

```
trainset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 6832 entries, 2092 to 426
Data columns (total 8 columns):
 #   Column                Non-Null Count  Dtype  
---  -
 0   datetime              6832 non-null  object 
 1   Count                 6832 non-null  int64  
 2   Day_of_Week           6832 non-null  object 
 3   Precipitation         6832 non-null  object 
 4   High_Winds            6832 non-null  bool    
 5   Road_Surface_Conditions 6832 non-null  object 
 6   Light_Conditions      6832 non-null  object 
 7   Urban_or_Rural_Area    6832 non-null  object 
dtypes: bool(1), int64(1), object(6)
memory usage: 433.7+ KB
```

From an overview, we can see that the datetime column datatype needs to be converted from an object to datetime64[ns]. In the non-null column, it looks like there are no missing values; however this is will be checked and confirmed.

## Convert datetime column

In [9]:

```
trainset['datetime'] = pd.to_datetime(trainset['datetime'])
trainset.head()
```

Out[9]:

	datetime	Count	Day_of_Week	Precipitation	High_Winds	Road_Surface_Conditions	Lig
2092	2018-04-01 05:00:00	2	Sunday	Fine	False	Wet or damp	D
6441	2018-10-03 18:00:00	18	Wednesday	Fine	False	Dry	
6172	2018-09-22 10:00:00	13	Saturday	Fine	False	Dry	
4187	2018-06-29 16:00:00	32	Friday	Fine	False	Dry	
2330	2018-04-11 10:00:00	12	Wednesday	Fine	False	Wet or damp	

In [10]:

```
#repeat for testset
testset['datetime'] = pd.to_datetime(testset['datetime'])
```

## Value counts in each column (categorical)

In [11]:

```
trainset['Day_of_Week'].value_counts()
```

Out[11]:

```
Sunday      995
Saturday    994
Monday      987
Friday      974
Wednesday   965
Thursday    963
Tuesday     954
Name: Day_of_Week, dtype: int64
```

In [12]:

```
trainset['Precipitation'].value_counts()
```

Out[12]:

```
Fine                5996
Raining             645
Snowing             92
Other                44
Unknown             35
Fog or mist         19
Data missing or out of range 1
Name: Precipitation, dtype: int64
```

80 records are classed as 'other', 'unknown' or 'data missing out of range'. Missing values will be dealt with in next step

In [13]:

```
trainset['High_Winds'].value_counts()
```

Out[13]:

```
False    6776
True       56
Name: High_Winds, dtype: int64
```

In [14]:

```
trainset['Road_Surface_Conditions'].value_counts()
```

Out[14]:

```
Dry                5056
Wet or damp        1631
Snow                90
Frost or ice        42
Data missing or out of range 8
Flood over 3cm. deep 5
Name: Road_Surface_Conditions, dtype: int64
```

8 records classed as 'data missing or out of range'

In [15]:

```
trainset['Light_Conditions'].value_counts()
```

Out[15]:

```
Daylight            3859
Darkness - lights lit 2453
Darkness - no lighting 414
Darkness - lighting unknown 84
Darkness - lights unlit 22
Name: Light_Conditions, dtype: int64
```



In [16]:

```
trainset['Urban_or_Rural_Area'].value_counts()
```

Out[16]:

```
Urban      5525  
Rural      1307  
Name: Urban_or_Rural_Area, dtype: int64
```

## Data Transformation

### Missing Values

The `Precipitation`, `Road_Surface_Conditions`, columns have records where the records are `Unknown`, Data missing or out of range and `Other`. Here we are going to replace the missing values using median.

### Precipitation column

In [17]:

```
#Filter for missing values
filter_list = ['Data missing or out of range', 'Unknown', 'Other']
trainset[trainset.Precipitation.isin(filter_list)]
```

Out[17]:

	datetime	Count	Day_of_Week	Precipitation	High_Winds	Road_Surface_Conditions	Lig
8451	2018-12-28 05:00:00	1	Friday	Other	False	Wet or damp	lig
1301	2018-02-26 02:00:00	1	Monday	Other	False	Dry	
5844	2018-09-08 08:00:00	5	Saturday	Unknown	False	Dry	
5399	2018-08-20 04:00:00	2	Monday	Unknown	False	Wet or damp	lig
7391	2018-11-13 01:00:00	2	Tuesday	Unknown	False	Wet or damp	
...	...	...	...	...	...	...	
27	2018-01-02 04:00:00	1	Tuesday	Other	False	Dry	D
8120	2018-12-14 04:00:00	1	Friday	Unknown	False	Data missing or out of range	D
2138	2018-04-03 04:00:00	1	Tuesday	Unknown	False	Data missing or out of range	
6047	2018-09-17 00:00:00	2	Monday	Unknown	False	Dry	
2368	2018-04-13 01:00:00	1	Friday	Other	False	Wet or damp	D

80 rows × 8 columns

We can see from the `trainset['Precipitation'].value_counts()` line, that all rows add up to 80 records; as evidenced in above table

In [18]:

```
# replace these rows with nan
trainset['Precipitation'] = trainset['Precipitation'].replace(['Other', 'Unknown', 'Data missing or out of range'], np.NaN)
```

In [19]:

```
#check results
filter_list = ['Data missing or out of range', 'Unknown', 'Other']
trainset[trainset.Precipitation.isin(filter_list)]
```

Out[19]:

datetime	Count	Day_of_Week	Precipitation	High_Winds	Road_Surface_Conditions	Light_C
----------	-------	-------------	---------------	------------	-------------------------	---------

In [20]:

```
#repeat for testset
testset['Precipitation'] = testset['Precipitation'].replace(['Other', 'Unknown',
'Data missing or out of range'], np.NaN)
```

## Road Surface Conditions column

In [21]:

```
#filter for missing values
filter_list = ['Data missing or out of range', 'Unknown', 'Other']
trainset[trainset.Road_Surface_Conditions.isin(filter_list)]
```

Out[21]:

	datetime	Count	Day_of_Week	Precipitation	High_Winds	Road_Surface_Conditions	Lig
1810	2018-03-20 03:00:00	1	Tuesday	NaN	False	Data missing or out of range	D
4973	2018-08-02 02:00:00	2	Thursday	NaN	False	Data missing or out of range	
4761	2018-07-24 02:00:00	1	Tuesday	Fine	False	Data missing or out of range	D
7553	2018-11-19 23:00:00	2	Monday	NaN	False	Data missing or out of range	
5753	2018-09-04 03:00:00	2	Tuesday	NaN	False	Data missing or out of range	
6687	2018-10-14 04:00:00	4	Sunday	Fine	False	Data missing or out of range	lig
8120	2018-12-14 04:00:00	1	Friday	NaN	False	Data missing or out of range	D
2138	2018-04-03 04:00:00	1	Tuesday	NaN	False	Data missing or out of range	

In [22]:

```
#replace these rows with nan
trainset['Road_Surface_Conditions'] = trainset['Road_Surface_Conditions'].replace(
(['Other', 'Unknown', 'Data missing or out of range'], np.NaN)
```

In [23]:

```
#check results
filter_list = ['Data missing or out of range', 'Unknown', 'Other']
trainset[trainset.Road_Surface_Conditions.isin(filter_list)]
```

Out[23]:

datetime	Count	Day_of_Week	Precipitation	High_Winds	Road_Surface_Conditions	Light_C
----------	-------	-------------	---------------	------------	-------------------------	---------

In [24]:

```
#repeat for testset
testset['Road_Surface_Conditions'] = testset['Road_Surface_Conditions'].replace(
(['Other', 'Unknown', 'Data missing or out of range'], np.NaN)
```

Check for missing values

In [25]:

```
trainset.isna().sum()
```

Out[25]:

```
datetime          0
Count             0
Day_of_Week       0
Precipitation     80
High_Winds        0
Road_Surface_Conditions  8
Light_Conditions  0
Urban_or_Rural_Area  0
dtype: int64
```

- Replace the missing values with mode in the dataset

In [26]:

```
from sklearn.impute import SimpleImputer

imputer = SimpleImputer(strategy="most_frequent")

trainset_categorical = trainset["Precipitation"].values
trainset_numerical = trainset.drop("Precipitation", axis=1)

imputer.fit(trainset_numerical)
```

Out[26]:

```
SimpleImputer(add_indicator=False, copy=True, fill_value=None,
              missing_values=nan, strategy='most_frequent', verbose=
0)
```

In [27]:

```
transformed = imputer.transform(trainset_numerical)
```

In [28]:

```
trainset = pd.DataFrame(transformed, columns=trainset_numerical.columns)

# add the categorical variable back
trainset['Precipitation'] = trainset_categorical

# check if there are missing values again
trainset.isnull().sum()
```

Out[28]:

```
datetime          0
Count             0
Day_of_Week       0
High_Winds        0
Road_Surface_Conditions  0
Light_Conditions  0
Urban_or_Rural_Area  0
Precipitation     80
dtype: int64
```

In [29]:

```
from sklearn.impute import SimpleImputer

imputer = SimpleImputer(strategy="most_frequent")

trainset_categorical = trainset["Road_Surface_Conditions"].values
trainset_numerical = trainset.drop("Road_Surface_Conditions", axis=1)

imputer.fit(trainset_numerical)
```

Out[29]:

```
SimpleImputer(add_indicator=False, copy=True, fill_value=None,
              missing_values=nan, strategy='most_frequent', verbose=
0)
```

In [30]:

```
transformed = imputer.transform(trainset_numerical)
```

In [31]:

```
trainset = pd.DataFrame(transformed, columns=trainset_numerical.columns)

# add the categorical variable back
trainset['Road_Surface_Conditions'] = trainset_categorical

# check if there are missing values again
trainset.isnull().sum()
```

Out[31]:

```
datetime          0
Count             0
Day_of_Week       0
High_Winds        0
Light_Conditions  0
Urban_or_Rural_Area  0
Precipitation     0
Road_Surface_Conditions  0
dtype: int64
```

## Transform the test dataset

In [32]:

```
testset.isnull().sum()
```

Out[32]:

```
datetime          0
Count             0
Day_of_Week       0
Precipitation     20
High_Winds        0
Road_Surface_Conditions  3
Light_Conditions  0
Urban_or_Rural_Area  0
dtype: int64
```

In [33]:

```
# precipitation
testset_categorical = testset["Precipitation"].values
testset_numerical = testset.drop("Precipitation", axis=1)

transformed = imputer.transform(testset_numerical)

testset = pd.DataFrame(transformed, columns=testset_numerical.columns)

# add the categorical variable back
testset['Precipitation'] = testset_categorical

# check if there are missing values again
testset.isnull().sum()
```

Out[33]:

```
datetime          0
Count             0
Day_of_Week       0
High_Winds        0
Road_Surface_Conditions  0
Light_Conditions  0
Urban_or_Rural_Area  0
Precipitation     20
dtype: int64
```

In [34]:

```
#road surface conditions
testset_categorical = testset["Road_Surface_Conditions"].values
testset_numerical = testset.drop("Road_Surface_Conditions", axis=1)

transformed = imputer.transform(testset_numerical)

testset = pd.DataFrame(transformed, columns=testset_numerical.columns)

# add the categorical variable back
testset['Road_Surface_Conditions'] = testset_categorical

# check if there are missing values again
testset.isnull().sum()
```

Out[34]:

```
datetime          0
Count             0
Day_of_Week       0
High_Winds        0
Light_Conditions  0
Urban_or_Rural_Area  0
Precipitation     0
Road_Surface_Conditions  0
dtype: int64
```

There are now no missing values in the train and test dataset

# Feature Engineering

- Adding month, hour of the day, year, season as extra columns.
- Mapping the weekdays to a dict to use the actual weekdays in words like 'Monday', etc. Python usually indicates Monday as 0, Tuesday as 1 and so on.
- Adding season based on the UK calendar

In [35]:

```
def season_calc(month):
    """adding season based on the UK weather """
    if month in [3,4,5]:
        return "spring"
    if month in [6,7,8]:
        return "summer"
    if month in [9,10,11]:
        return "autumn"
    else:
        return "winter"
```

In [36]:

```
trainset['year'] = trainset.datetime.dt.year
trainset['month'] = trainset.datetime.dt.month
trainset['day'] = trainset.datetime.dt.day
trainset['hour'] = trainset.datetime.dt.hour
trainset['season'] = trainset.datetime.dt.month.apply(season_calc)
```

In [37]:

```
#check the results
trainset.head()
```

Out[37]:

	datetime	Count	Day_of_Week	High_Winds	Light_Conditions	Urban_or_Rural_Area	Precipi
0	2018-04-01 05:00:00	2	Sunday	False	Darkness - lights lit	Urban	
1	2018-10-03 18:00:00	18	Wednesday	False	Daylight	Urban	
2	2018-09-22 10:00:00	13	Saturday	False	Daylight	Rural	
3	2018-06-29 16:00:00	32	Friday	False	Daylight	Urban	
4	2018-04-11 10:00:00	12	Wednesday	False	Daylight	Urban	



- Repeat steps for test data set

In [38]:

```
testset['year'] = testset.datetime.dt.year
testset['month'] = testset.datetime.dt.month
testset['day'] = testset.datetime.dt.day
testset['hour'] = testset.datetime.dt.hour
testset['season'] = testset.datetime.dt.month.apply(season_calc)
```

## Create dummy variables

As predictive models will be build with the effect of area, dummy variables need to be created as these are categorical variables

In [39]:

```
from sklearn.preprocessing import OneHotEncoder

one_hot_encoder = OneHotEncoder(drop="first", sparse=False)

# the input to the encoder must be a 2-d numpy array,
# so we take the column, extract their values and reshape the array to be 2-d
cat_vals = trainset['Urban_or_Rural_Area'].values.reshape(-1,1)

transformed = one_hot_encoder.fit_transform(cat_vals)

# put the transformed data as columns in the dataframe
col_names = one_hot_encoder.categories_[0].tolist()[1:]
for i, col_name in enumerate(col_names):
    trainset[col_name] = transformed[:,i]

# check if the dummies are produced correctly
trainset.head(2)
```

Out[39]:

	datetime	Count	Day_of_Week	High_Winds	Light_Conditions	Urban_or_Rural_Area	Precipi
0	2018-04-01 05:00:00	2	Sunday	False	Darkness - lights lit	Urban	
1	2018-10-03 18:00:00	18	Wednesday	False	Daylight	Urban	

In [40]:

```
# delete the categorical column
del trainset['Urban_or_Rural_Area']
```

Create dummy variables in the test set

In [41]:

```
cat_vals = testset['Urban_or_Rural_Area'].values.reshape(-1,1)
transformed = one_hot_encoder.transform(cat_vals)

for i, col_name in enumerate(col_names):
    testset[col_name] = transformed[:,i]

# check if the dummies are produced correctly
testset.head()
```

Out[41]:

	datetime	Count	Day_of_Week	High_Winds	Light_Conditions	Urban_or_Rural_Area	Precipi
0	2018-03-06 16:00:00	26	Tuesday	False	Daylight	Urban	
1	2018-05-20 17:00:00	23	Sunday	False	Daylight	Urban	
2	2018-04-03 14:00:00	24	Tuesday	False	Daylight	Urban	
3	2018-01-31 18:00:00	32	Wednesday	False	Darkness - lights lit	Urban	
4	2018-04-26 11:00:00	19	Thursday	False	Daylight	Urban	

In [42]:

```
# delete the categorical column
del testset['Urban_or_Rural_Area']
```

## Finalise trainset for analysis

In [43]:

```
trainset.head(2)
```

Out[43]:

	datetime	Count	Day_of_Week	High_Winds	Light_Conditions	Precipitation	Road_Surface_
0	2018-04-01 05:00:00	2	Sunday	False	Darkness - lights lit	Fine	W
1	2018-10-03 18:00:00	18	Wednesday	False	Daylight	Fine	

- Sort series data in line with date

In [44]:

```
trainset = trainset.sort_values(by=['datetime'])
trainset.head()
```

Out[44]:

	datetime	Count	Day_of_Week	High_Winds	Light_Conditions	Precipitation	Road_Surfa
824	2018-01-01 00:00:00	18	Monday	False	Darkness - lights lit		Fine
263	2018-01-01 02:00:00	14	Monday	False	Darkness - lights lit		Fine
3993	2018-01-01 03:00:00	10	Monday	False	Darkness - lights lit		Fine
6325	2018-01-01 04:00:00	8	Monday	False	Darkness - lights lit		Fine
3709	2018-01-01 05:00:00	9	Monday	False	Darkness - lights lit		Fine

In [45]:

```
trainset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 6832 entries, 824 to 2458
Data columns (total 13 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   datetime                             6832 non-null  datetime64[ns]
1   Count                                6832 non-null  object
2   Day_of_Week                          6832 non-null  object
3   High_Winds                           6832 non-null  object
4   Light_Conditions                     6832 non-null  object
5   Precipitation                        6832 non-null  object
6   Road_Surface_Conditions              6832 non-null  object
7   year                                 6832 non-null  int64
8   month                                6832 non-null  int64
9   day                                  6832 non-null  int64
10  hour                                 6832 non-null  int64
11  season                               6832 non-null  object
12  Urban                                6832 non-null  float64
dtypes: datetime64[ns](1), float64(1), int64(4), object(7)
memory usage: 747.2+ KB
```

- Reinstate int64 to Count column

In [46]:

```
trainset['Count'] = trainset['Count'].astype(str).astype(int)
```

In [47]:

```
trainset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 6832 entries, 824 to 2458
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	datetime	6832 non-null	datetime64[ns]
1	Count	6832 non-null	int64
2	Day_of_Week	6832 non-null	object
3	High_Winds	6832 non-null	object
4	Light_Conditions	6832 non-null	object
5	Precipitation	6832 non-null	object
6	Road_Surface_Conditions	6832 non-null	object
7	year	6832 non-null	int64
8	month	6832 non-null	int64
9	day	6832 non-null	int64
10	hour	6832 non-null	int64
11	season	6832 non-null	object
12	Urban	6832 non-null	float64

```
dtypes: datetime64[ns](1), float64(1), int64(5), object(6)
```

```
memory usage: 747.2+ KB
```

- Convert Day\_of\_Week and season to category

In [48]:

```
for col in ['spring', 'summer', 'autumn', 'winter']:
    trainset['season'] = trainset['season'].astype('category')
```

In [49]:

```
for col in ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']:
    trainset['Day_of_Week'] = trainset['Day_of_Week'].astype('category')
```

In [50]:

```
trainset.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 6832 entries, 824 to 2458
```

```
Data columns (total 13 columns):
```

#	Column	Non-Null Count	Dtype
0	datetime	6832 non-null	datetime64[ns]
1	Count	6832 non-null	int64
2	Day_of_Week	6832 non-null	category
3	High_Winds	6832 non-null	object
4	Light_Conditions	6832 non-null	object
5	Precipitation	6832 non-null	object
6	Road_Surface_Conditions	6832 non-null	object
7	year	6832 non-null	int64
8	month	6832 non-null	int64
9	day	6832 non-null	int64
10	hour	6832 non-null	int64
11	season	6832 non-null	category
12	Urban	6832 non-null	float64

```
dtypes: category(2), datetime64[ns](1), float64(1), int64(5), object
(4)
```

```
memory usage: 654.4+ KB
```

- Select columns for analysis

In [51]:

```
trainset = trainset.drop(columns=['High_Winds', 'Light_Conditions', 'Precipitation', 'Road_Surface_Conditions'])
trainset.head(2)
```

Out[51]:

	datetime	Count	Day_of_Week	year	month	day	hour	season	Urban
<b>824</b>	2018-01-01 00:00:00	18	Monday	2018	1	1	0	winter	1.0
<b>263</b>	2018-01-01 02:00:00	14	Monday	2018	1	1	2	winter	1.0

In [52]:

```
#set datetime as index
trainset = trainset.set_index('datetime')
trainset.index
```

Out[52]:

```
DatetimeIndex(['2018-01-01 00:00:00', '2018-01-01 02:00:00',
               '2018-01-01 03:00:00', '2018-01-01 04:00:00',
               '2018-01-01 05:00:00', '2018-01-01 06:00:00',
               '2018-01-01 07:00:00', '2018-01-01 09:00:00',
               '2018-01-01 10:00:00', '2018-01-01 11:00:00',
               ...,
               '2018-12-31 08:00:00', '2018-12-31 09:00:00',
               '2018-12-31 11:00:00', '2018-12-31 15:00:00',
               '2018-12-31 16:00:00', '2018-12-31 17:00:00',
               '2018-12-31 18:00:00', '2018-12-31 20:00:00',
               '2018-12-31 22:00:00', '2018-12-31 23:00:00'],
              dtype='datetime64[ns]', name='datetime', length=6832,
              freq=None)
```

In [53]:

```
trainset.head(3)
```

Out[53]:

	Count	Day_of_Week	year	month	day	hour	season	Urban
<b>datetime</b>								
<b>2018-01-01 00:00:00</b>	18	Monday	2018	1	1	0	winter	1.0
<b>2018-01-01 02:00:00</b>	14	Monday	2018	1	1	2	winter	1.0
<b>2018-01-01 03:00:00</b>	10	Monday	2018	1	1	3	winter	1.0

## Repeat for test dataset

In [54]:

```
testset = testset.sort_values(by=['datetime'])
```

In [55]:

```
testset['Count'] = testset['Count'].astype(str).astype(int)
```

In [56]:

```
for col in ['spring', 'summer', 'autumn', 'winter']:
    trainset['season'] = trainset['season'].astype('category')
```

In [57]:

```
for col in ['Monday', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']:
    trainset['Day_of_Week'] = trainset['Day_of_Week'].astype('category')
```

In [58]:

```
testset = testset.drop(columns=['High_Winds', 'Light_Conditions', 'Precipitation', 'Road_Surface_Conditions'])
```

In [59]:

```
testset = testset.set_index('datetime')
```

## Visualisations

### Heatmap


Plotting a map of hourly vs weekdays number of accidents to variation across any particular week

In [60]:

```
hour_weekday = trainset.pivot_table(values='Count', index='hour', columns = 'Day_of_Week', aggfunc = 'mean')
```

In [61]:

```
#plotting a heatmap with a colorbar; the colorbar shows the number of accidents
_ = plt.figure(figsize=(12, 8))
ax = sns.heatmap(hour_weekday.sort_index(ascending = False), cmap='viridis')
#_ = plot title
_ = ax.set_title("Average number of accidents for each hour of each weekday averaged over 2018", fontsize = 14)
```




- The heatmap above shows that the average number of accidents from Monday to Friday is below 10 in the night. This increases between 7:30am and 8:30am, slows down slightly around midday and picks up again between 3pm and 6pm.
- Over the weekend, the average number of accidents rises above 12 in the night to early morning. Overall accidents occur less on the weekends.
- During the weekdays, the reason for the higher average number of accidents occurring between 7:30am -8:30am and 3pm - 6pm can be explained by peak times where individuals are rushing to work and school.

## Average Hourly 2018

### Plotting average hourly accidents over 2018

In [62]:

```
#Plotting average hourly number of accidents observed over the entire period
trainset.groupby('hour')['Count'].mean().plot(figsize = (10,5))
_ = plt.ylabel('Number of Accidents')
_ = plt.ylim([0, max(trainset.groupby('hour')['Count'].mean()) + 30])
_ = plt.xticks(trainset['hour'].unique())
_ = plt.title('Hourly Accident Occurences in 2018')
```



- As expected the hourly accidents peak in the early hours in the day and 3pm-5:30pm approximately due to rush hour/peak time. This was also shown in the heatmap plotted above

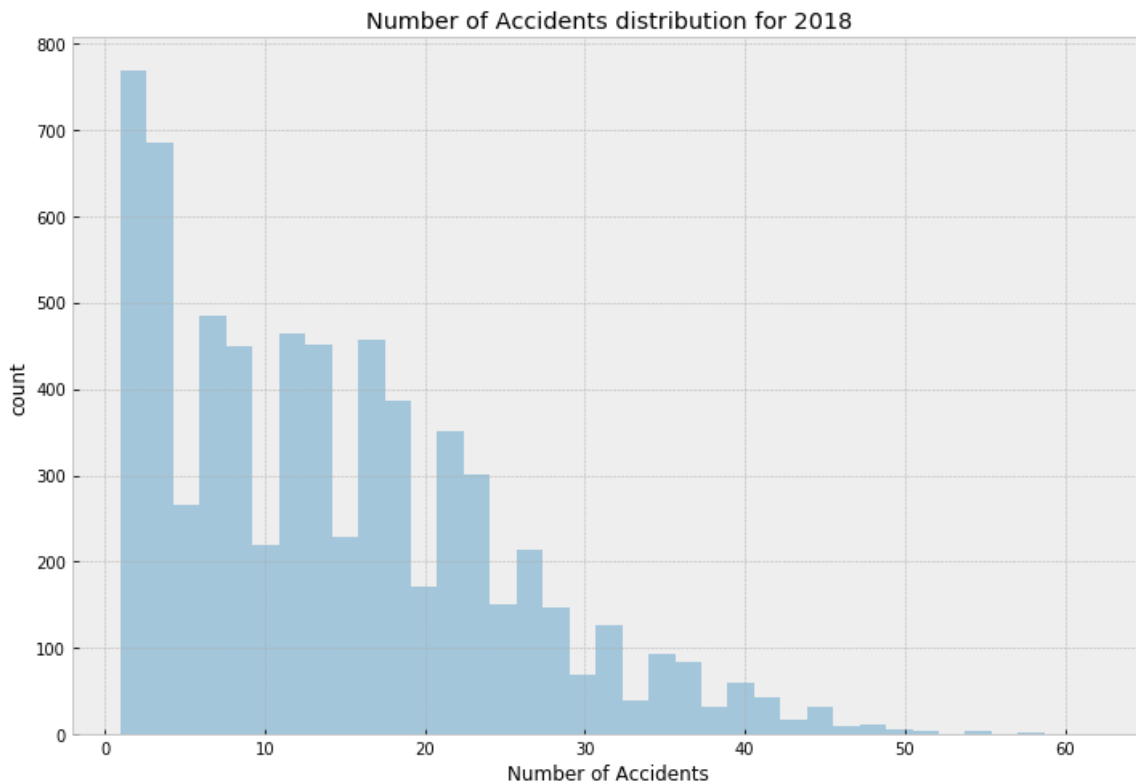
## Plotting a histogram

### Visualise the distribution of the number of accidents recorded



In [63]:

```
_ = plt.figure(figsize = (12,8))  
_ = sns.distplot(trainset['Count'], kde=False)  
_ = plt.title('Number of Accidents distribution for 2018')  
_ = plt.xlabel('Number of Accidents')  
_ = plt.ylabel('count')
```




The histogram above shows:

- for most of the days in 2018, there were around 5 - 25 accidents a day
- highest number of accidents in a day was around 58 accidents
- There are no outliers/anomalies in the data to deal with

## Average Monthly 2018

In [64]:

```
avm = trainset['Count'].resample('MS').mean()
avm.plot(figsize=(15, 6))
plt.show()
```



- The graph above shows the peaks in the number of accidents that occur during the summer and autumn months

## Time Series


Plot time series dataset

In [65]:

```
tsp = trainset['Count']
tsp.plot(figsize=(16,3))
```

Out[65]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f916a69e950>



# Correlations

As a predictive model will be built with questions on if area has an effect on the number of accidents, would be helpful to calculate correlation coefficient

In [66]:

```
# Let's calculate the pearsonr coefficient between the number of accidents and a
rea being urban or rural
scipy.stats.pearsonr(trainset['Count'], trainset['Urban'])
```

Out[66]:

```
(0.3072922466819525, 2.4416961710152814e-149)
```

- There is a positive correlation between the number of accidents and the area in which the accidents occurred.

# Time Series Model

To build time series model, additional preprocessing and transformation steps are required

## Additional Preprocessing and Transformation

In [67]:

```
#create a simple time series dataframe
tsd = pd.DataFrame(trainset["Count"], columns=['Count'])
```

In [68]:

```
tsd.head(3)
```

Out[68]:

	Count
datetime	
2018-01-01 00:00:00	18
2018-01-01 02:00:00	14
2018-01-01 03:00:00	10

In [69]:


```
#repeat for test data
test2 = pd.DataFrame(testset["Count"], columns=['Count'])
```

In [70]:

```
# Plot the time series  
tsd.plot(figsize=(16,3))
```

Out[70]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f9169baa810>



## Decompose time series

Decompose the time series data by `trend` , `seasonality` and `residuals`


In [71]:

```
from statsmodels.tsa.seasonal import seasonal_decompose

# specify the number of observations in a cycle
decomposition = seasonal_decompose(tsd['Count'], freq=365)

trend = decomposition.trend
seasonal = decomposition.seasonal
residual = decomposition.resid

plt.figure(figsize=(12, 8))
plt.subplot(411)
plt.plot(tsd['Count'], label='Original')
plt.legend(loc='best')
plt.subplot(412)
plt.plot(trend, label='Trend')
plt.legend(loc='best')
plt.subplot(413)
plt.plot(seasonal, label='Seasonality')
plt.legend(loc='best')
plt.subplot(414)
plt.plot(residual, label='Residuals')
plt.legend(loc='best')
plt.tight_layout()
```



- There seems to be seasonality during with the year, with an upward trend in the number of accidents in the summer and autumn months.
- The hourly change in the number of accidents has no cyclic behaviour.
- There does not seem to be any random fluctuations in the series.

### Testing for stationarity

Use ADF and KPSS test totest for stationarity

In [72]:

```
from statsmodels.tsa.stattools import adfuller, kpss

for x in ["Count"]:
    adf_pval = adfuller(tsd[x])[1]
    print(x)
    print(f"ADF, p-value: {adf_pval}")
    kpss_stat, kpss_pval, lags, crit_vals = kpss(trainset['Count'])
    print(f"KPSS, p-value: {kpss_pval}")
```

Count

ADF, p-value: 1.2211447254837683e-22

KPSS, p-value: 0.01811627239901272

- Based on the p-values of the above tests; the ADF test fails to reject the null of a unit root, and the KPSS test rejects the null of stationarity. The next steps will be done to stationarize the data.

## Log Transformation

In [73]:

```
for x in ["Count"]:
    print(x)
    logs = np.log(tsd[x])
    adf_pval = adfuller(logs)[1]
    print(f"ADF, p-value: {adf_pval}")
    kpss_stat, kpss_pval, lags, crit_vals = kpss(logs)
    print(f"KPSS, p-value: {kpss_pval}")
```

Count

ADF, p-value: 2.465969207267234e-27

KPSS, p-value: 0.03486597650163498

- After log transformation, the ADF still fails to reject the null of a unit root and the KPSS test rejects the null of stationarity. Next step will be to perform differencing

## Differencing

In [74]:

```
tsd_diff = tsd.diff()
tsd_diff.head(3)
```

Out[74]:

	Count
datetime	
2018-01-01 00:00:00	NaN
2018-01-01 02:00:00	-4.0
2018-01-01 03:00:00	-4.0

In [75]:

```
#drop the first missing value differencing created
tsd_diff.dropna(inplace=True)
```

In [76]:

```
#conduct adf and kpss test
for x in ["Count"]:
    print(x)
    adf_pval = adfuller(tsd_diff[x])[1]
    print("ADF, p-value:", adf_pval)
    kpss_stat, kpss_pval, lags, crit_vals = kpss(tsd_diff[x])
    print("KPSS, p-value:", kpss_pval)
```

Count

ADF, p-value: 0.0

KPSS, p-value: 0.1

```
/opt/anaconda3/lib/python3.7/site-packages/statsmodels/tsa/stattools.py:1710: InterpolationWarning: p-value is greater than the indicated p-value
  warn("p-value is greater than the indicated p-value", InterpolationWarning)
```


As a result of the first differencing, both tests indicate that the series has become stationary; the ADF rejects the null of unit root and the KPSS test fails to reject the null of stationarity at 0.05 significance level.

In [77]:

```
#plot the differenced data
plt.subplot(311)
tsd_diff['Count'].plot(figsize=(16, 6), legend=True)
```

Out[77]:

&lt;matplotlib.axes.\_subplots.AxesSubplot at 0x7f916a09bdd0&gt;



In [78]:

```
train_diff = tsd_diff
```

- Apply differencing to test data

In [79]:

```
test_diff = test2.diff().dropna()
```

In [80]:

```
testdiff = testset[['Count', 'Urban']].shift().replace(np.NaN, 0.0)
```



## Baseline

Will use a persistence baseline for this analysis as it is most common baseline method used for supervised machine learning; and can also be used on time series data.

In [81]:

```
test_diff['Count'].head()
```

Out[81]:

```
datetime
2018-01-01 08:00:00    -1.0
2018-01-01 15:00:00     9.0
2018-01-01 22:00:00   -12.0
2018-01-01 23:00:00    3.0
2018-01-02 05:00:00   -5.0
Name: Count, dtype: float64
```

In [82]:

```
test_diff['Count'].shift().head()
```

Out[82]:

```
datetime
2018-01-01 08:00:00     NaN
2018-01-01 15:00:00    -1.0
2018-01-01 22:00:00     9.0
2018-01-01 23:00:00   -12.0
2018-01-02 05:00:00     3.0
Name: Count, dtype: float64
```

In [83]:

```
mse = mean_squared_error(test_diff['Count'][1:], test_diff['Count'].shift()[1:])
np.sqrt(mse)
```

Out[83]:

```
20.72135517709077
```

In [84]:

```
mda = get_mda(test_diff['Count'][1:], test_diff['Count'].shift()[1:])
mda
```

Out[84]:

```
0.318475073313783
```

## ARIMA

### Determine the parameters

The most common way to determine the  $p$  and  $q$  parameters is by using ACF and PACF plots.

In [85]:

```
from statsmodels.tsa.stattools import acf, pacf


lag_acf = acf(train_diff['Count'], nlags=10)
lag_pacf = pacf(train_diff['Count'], nlags=10, method='ols')
```

In [86]:

```
plt.figure(figsize=(12, 4))

plt.subplot(121)
plt.plot(lag_acf)
plt.axhline(y=0, linestyle='--', color='gray')
plt.axhline(y=-1.96/np.sqrt(len(train_diff)), linestyle='--', color='gray')
plt.axhline(y=1.96/np.sqrt(len(train_diff)), linestyle='--', color='gray')
plt.title('Autocorrelation Function')

plt.subplot(122)
plt.plot(lag_pacf)
plt.axhline(y=0, linestyle='--', color='gray')
plt.axhline(y=-1.96/np.sqrt(len(train_diff)), linestyle='--', color='gray')
plt.axhline(y=1.96/np.sqrt(len(train_diff)), linestyle='--', color='gray')
plt.title('Partial Autocorrelation Function')
plt.tight_layout()
```



From results the ACF and PACF graphs above, set  $q$  as 1 and  $p$  as 1. As the lagged correlation is not negative, can conclude the series has not been overdifferenced.

## Hyperparameter Tuning

Here the ARIMA hyperparameters will be tuned using a grid search

In [87]:

```
#confirm using code to find best parameters
resDiff = sm.tsa.arma_order_select_ic(train_diff, max_ar=1, max_ma=1, ic='aic',
trend='c')
print('ARMA(p,q) =', resDiff['aic_min_order'], 'is the best.')
```

ARMA(p,q) = (1, 1) is the best.

Therefore ARIMA model order would be (1,0,1) as we are using differenced series

## Build ARIMA Model

In [88]:

```
endog = train_diff['Count']
endog.head(3)
```

Out[88]:

```
datetime
2018-01-01 02:00:00    -4.0
2018-01-01 03:00:00    -4.0
2018-01-01 04:00:00    -2.0
Name: Count, dtype: float64
```

In [89]:

```
from statsmodels.tsa.arma_model import ARIMA
```

In [90]:

```
arima = ARIMA(endog, order=(0, 0, 1)).fit(solver="bfgs", disp=0)

# print the significance of the variables
print(arima.summary().tables[1])
```

```
=====
=====
              coef      std err          z      P>|z|      [0.025
0.975]
-----
-----
const          -0.0011      0.087      -0.013      0.990      -0.171
0.169
ma.L1.Count     -0.1170      0.013     -8.779      0.000      -0.143
-0.091
=====
=====
```

```
/opt/anaconda3/lib/python3.7/site-packages/statsmodels/tsa/base/tsa_
model.py:219: ValueWarning: A date index has been provided, but it h
as no associated frequency information and so will be ignored when
e.g. forecasting.
  ' ignored when e.g. forecasting.', ValueWarning)
```

## Evaluate ARIMA on test data

- Retrain new model for each instance of the test data

In [91]:

```
# buffers keeping previously seen endogenous variables
history_endog = [x for x in endog]
predictions = []

# for each test observation, take the first 300 for convenience
for i, test_obs in enumerate(test_diff['Count'][:200]):

    # build a model using the current buffers for endogenous variables
    model = ARIMA(history_endog, order=(0, 0, 1)).fit(solver="bfgs", disp=0)

    # forecast the value for the test instance, supplying corresponding exogenous variables
    yhat = model.forecast()[0][0]

    # remember the forecasted value
    predictions.append(yhat)

    # update the buffers for the endogenous variables
    history_endog.append(test_obs)

print('predicted=%f, expected=%f' % (yhat, test_obs))
```

predicted=0.974270, expected=-1.000000  
predicted=0.229527, expected=9.000000  
predicted=-1.026657, expected=-12.000000  
predicted=1.285117, expected=3.000000  
predicted=-0.202581, expected=-5.000000  
predicted=0.560882, expected=6.000000  
predicted=-0.639906, expected=4.000000  
predicted=-0.545189, expected=-4.000000  
predicted=0.404364, expected=9.000000  
predicted=-1.009711, expected=-1.000000  
predicted=-0.001218, expected=1.000000  
predicted=-0.117548, expected=3.000000  
predicted=-0.365711, expected=-7.000000  
predicted=0.779157, expected=-9.000000  
predicted=1.145724, expected=-2.000000  
predicted=0.366670, expected=5.000000  
predicted=-0.544973, expected=-7.000000  
predicted=0.755350, expected=28.000000  
predicted=-3.208414, expected=-10.000000  
predicted=0.802281, expected=-6.000000  
predicted=0.803712, expected=-3.000000  
predicted=0.448362, expected=2.000000  
predicted=-0.184201, expected=-7.000000  
predicted=0.803924, expected=17.000000  
predicted=-1.918527, expected=-2.000000  
predicted=0.010010, expected=-3.000000  
predicted=0.356694, expected=-8.000000  
predicted=0.988712, expected=0.000000  
predicted=0.115837, expected=-1.000000  
predicted=0.130766, expected=7.000000  
predicted=-0.814102, expected=9.000000  
predicted=-1.159763, expected=-6.000000  
predicted=0.573087, expected=3.000000  
predicted=-0.286938, expected=1.000000  
predicted=-0.151768, expected=-10.000000  
predicted=1.165995, expected=-3.000000  
predicted=0.491766, expected=-8.000000  
predicted=1.001662, expected=1.000000  
predicted=-0.002063, expected=-1.000000  
predicted=0.115596, expected=1.000000  
predicted=-0.106838, expected=0.000000  
predicted=-0.014892, expected=17.000000  
predicted=-2.011617, expected=-8.000000  
predicted=0.708203, expected=22.000000  
predicted=-2.525838, expected=-17.000000  
predicted=1.728126, expected=-15.000000  
predicted=1.987659, expected=15.000000  
predicted=-1.555695, expected=27.000000  
predicted=-3.388006, expected=-10.000000  
predicted=0.789454, expected=4.000000  
predicted=-0.380408, expected=-31.000000  
predicted=3.660418, expected=10.000000  
predicted=-0.759836, expected=1.000000  
predicted=-0.211434, expected=-5.000000  
predicted=0.574442, expected=10.000000  
predicted=-1.132660, expected=-6.000000  
predicted=0.585490, expected=17.000000  
predicted=-1.976066, expected=2.000000  
predicted=-0.475685, expected=-20.000000  
predicted=2.355323, expected=5.000000  
predicted=-0.318829, expected=-4.000000

predicted=0.444474, expected=0.000000  
predicted=0.053470, expected=34.000000  
predicted=-4.098659, expected=-41.000000  
predicted=4.568136, expected=-3.000000  
predicted=0.930727, expected=-1.000000  
predicted=0.236528, expected=6.000000  
predicted=-0.712949, expected=-3.000000  
predicted=0.281222, expected=-4.000000  
predicted=0.527006, expected=4.000000  
predicted=-0.430605, expected=-3.000000  
predicted=0.315802, expected=6.000000  
predicted=-0.703781, expected=13.000000  
predicted=-1.690700, expected=-6.000000  
predicted=0.532485, expected=6.000000  
predicted=-0.675110, expected=-8.000000  
predicted=0.906206, expected=11.000000  
predicted=-1.249439, expected=-16.000000  
predicted=1.832086, expected=-8.000000  
predicted=1.215910, expected=20.000000  
predicted=-2.337307, expected=-11.000000  
predicted=1.080442, expected=26.000000  
predicted=-3.120186, expected=-29.000000  
predicted=3.280231, expected=-4.000000  
predicted=0.918267, expected=1.000000  
predicted=-0.012003, expected=18.000000  
predicted=-2.279598, expected=-19.000000  
predicted=2.126596, expected=-1.000000  
predicted=0.395795, expected=-3.000000  
predicted=0.429660, expected=5.000000  
predicted=-0.583369, expected=-3.000000  
predicted=0.305557, expected=1.000000  
predicted=-0.090358, expected=-2.000000  
predicted=0.240892, expected=8.000000  
predicted=-0.989236, expected=7.000000  
predicted=-1.016128, expected=-1.000000  
predicted=-0.002265, expected=2.000000  
predicted=-0.254594, expected=3.000000  
predicted=-0.413392, expected=-14.000000  
predicted=1.727986, expected=-6.000000  
predicted=0.978865, expected=17.000000  
predicted=-2.040987, expected=9.000000  
predicted=-1.399984, expected=-22.000000  
predicted=2.627645, expected=0.000000  
predicted=0.333084, expected=-4.000000  
predicted=0.550098, expected=21.000000  
predicted=-2.612239, expected=7.000000  
predicted=-1.220486, expected=-21.000000  
predicted=2.527647, expected=-7.000000  
predicted=1.210870, expected=2.000000  
predicted=-0.102646, expected=29.000000  
predicted=-3.709415, expected=-7.000000  
predicted=0.420631, expected=7.000000  
predicted=-0.838606, expected=-20.000000  
predicted=2.453730, expected=-8.000000  
predicted=1.330847, expected=-1.000000  
predicted=0.295036, expected=15.000000  
predicted=-1.877405, expected=4.000000  
predicted=-0.748098, expected=15.000000  
predicted=-1.999142, expected=-26.000000  
predicted=3.073732, expected=23.000000  
predicted=-2.571752, expected=-31.000000

```
predicted=3.708271, expected=2.000000
predicted=0.220734, expected=9.000000
predicted=-1.147138, expected=5.000000
predicted=-0.801762, expected=0.000000
predicted=-0.104378, expected=0.000000
predicted=-0.013399, expected=-1.000000
predicted=0.128797, expected=1.000000
predicted=-0.113448, expected=-15.000000
predicted=1.940798, expected=13.000000
predicted=-1.447229, expected=-2.000000
predicted=0.072013, expected=-9.000000
predicted=1.186411, expected=1.000000
predicted=0.022902, expected=-6.000000
predicted=0.786230, expected=26.000000
predicted=-3.309916, expected=-24.000000
predicted=2.738379, expected=26.000000
predicted=-3.106250, expected=11.000000
predicted=-1.869987, expected=-26.000000
predicted=3.230644, expected=-2.000000
predicted=0.697725, expected=8.000000
predicted=-0.976971, expected=12.000000
predicted=-1.732146, expected=-22.000000
predicted=2.721197, expected=-8.000000
predicted=1.431739, expected=13.000000
predicted=-1.552861, expected=15.000000
predicted=-2.213116, expected=0.000000
predicted=-0.293936, expected=-2.000000
predicted=0.229700, expected=-11.000000
predicted=1.501649, expected=-12.000000
predicted=1.797742, expected=27.000000
predicted=-3.380953, expected=-6.000000
predicted=0.352923, expected=13.000000
predicted=-1.697389, expected=-35.000000
predicted=4.510581, expected=8.000000
predicted=-0.473575, expected=-6.000000
predicted=0.749022, expected=-1.000000
predicted=0.235726, expected=-2.000000
predicted=0.301520, expected=9.000000
predicted=-1.182441, expected=-8.000000
predicted=0.925074, expected=6.000000
predicted=-0.691560, expected=-9.000000
predicted=1.128903, expected=2.000000
predicted=-0.120677, expected=5.000000
predicted=-0.698687, expected=14.000000
predicted=-1.998092, expected=12.000000
predicted=-1.892271, expected=-27.000000
predicted=3.422865, expected=-6.000000
predicted=1.276661, expected=0.000000
predicted=0.171050, expected=5.000000
predicted=-0.657949, expected=9.000000
predicted=-1.311879, expected=-2.000000
predicted=0.092800, expected=12.000000
predicted=-1.615931, expected=-24.000000
predicted=3.051851, expected=8.000000
predicted=-0.677168, expected=37.000000
predicted=-5.140886, expected=-35.000000
predicted=4.148415, expected=-5.000000
predicted=1.267164, expected=-3.000000
predicted=0.589842, expected=8.000000
predicted=-1.029960, expected=2.000000
predicted=-0.421251, expected=6.000000
```

```

predicted=-0.890786, expected=-15.000000
predicted=1.959286, expected=0.000000
predicted=0.270211, expected=-1.000000
predicted=0.174417, expected=17.000000
predicted=-2.338674, expected=0.000000
predicted=-0.324279, expected=-2.000000
predicted=0.232889, expected=6.000000
predicted=-0.800424, expected=-19.000000
predicted=2.531722, expected=-1.000000
predicted=0.488819, expected=26.000000
predicted=-3.555450, expected=-6.000000
predicted=0.341675, expected=-10.000000
predicted=1.443553, expected=-6.000000
predicted=1.036288, expected=-1.000000
predicted=0.282299, expected=18.000000
predicted=-2.471198, expected=-11.000000
predicted=1.191950, expected=-12.000000

```

In [92]:


```

#plot predictions on a graph to compare with observed values
pd.DataFrame({"true": test_diff['Count'][:200], "predicted": predictions}).plot(
    figsize=(16, 4))

```

Out[92]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f91518dc6d0>



- Above is a plot of the predicted values vs the observed models using ARIMA. The predicted values seems to match the peaks and troughs.

In [93]:

```

mse = mean_squared_error(test_diff['Count'][:200], predictions)
np.sqrt(mse)

```

Out[93]:

12.520498573775068

In [94]:

```

mda = get_mda(test_diff['Count'][:200], predictions)
mda

```

Out[94]:

0.7085427135678392



Here the ARIMA hyperparameters will be tuned using a grid search

In [95]:

```
from statsmodels.tsa.arima_model import ARIMA
```

## Further Models

### ARIMAX

#### Add exogenous variables

In [96]:

```
exog = trainset[['Count', 'Urban']].shift().replace(np.NaN, 0.0)  
exog.head(3)
```

Out[96]:

	Count	Urban
datetime		
2018-01-01 00:00:00	0.0	0.0
2018-01-01 02:00:00	18.0	1.0
2018-01-01 03:00:00	14.0	1.0

In [97]:

```
exogd = exog.diff().dropna()
exogd
```

Out[97]:

	Count	Urban
datetime		
2018-01-01 02:00:00	18.0	1.0
2018-01-01 03:00:00	-4.0	0.0
2018-01-01 04:00:00	-4.0	0.0
2018-01-01 05:00:00	-2.0	0.0
2018-01-01 06:00:00	1.0	0.0
...	...	...
2018-12-31 17:00:00	10.0	0.0
2018-12-31 18:00:00	1.0	0.0
2018-12-31 20:00:00	-7.0	0.0
2018-12-31 22:00:00	-2.0	0.0
2018-12-31 23:00:00	6.0	0.0

6831 rows × 2 columns

In [98]:

```
#rename exog
exog = exogd
```

In [99]:

```
endog = train_diff['Count']
endog
```

Out[99]:

```
datetime
2018-01-01 02:00:00    -4.0
2018-01-01 03:00:00    -4.0
2018-01-01 04:00:00    -2.0
2018-01-01 05:00:00     1.0
2018-01-01 06:00:00    -4.0
...
2018-12-31 17:00:00     1.0
2018-12-31 18:00:00    -7.0
2018-12-31 20:00:00    -2.0
2018-12-31 22:00:00     6.0
2018-12-31 23:00:00    -9.0
Name: Count, Length: 6831, dtype: float64
```

## Build ARIMAX model

In [100]:

```

from statsmodels.tsa.arima_model import ARIMA

arima = ARIMA(endog, exog=exog, order=(0, 0, 1)).fit(solver="bfgs", disp=0)

# print the significance of the variables
print(arima.summary().tables[1])

```

/opt/anaconda3/lib/python3.7/site-packages/statsmodels/tsa/base/tsa\_model.py:219: ValueWarning: A date index has been provided, but it has no associated frequency information and so will be ignored when e.g. forecasting.

' ignored when e.g. forecasting.', ValueWarning)

```

=====
=====

```

	coef	std err	z	P> z	[0.025
0.975]					
-----					
const	5.996e-05	4.61e-05	1.301	0.193	-3.04e-05
0.000					
Count	0.7021	0.009	77.038	0.000	0.684
0.720					
Urban	-0.4617	0.243	-1.902	0.057	-0.937
0.014					
ma.L1.Count	-1.0000	0.000	-2149.492	0.000	-1.001
-0.999					
=====					
=====					

In [ ]:

**Evaluate model on test data**

In [101]:

```
# buffers keeping previously seen endogenous and exogenous variables
history_endog = [x for x in endog]
history_exog = [x for x in exog.values]
predictions = []

# create exogenous variables also for the test data
test_exog = testdiff[['Count', 'Urban']].shift().replace(np.NaN, 0.0)

# for each test observation, take the next 300 for convenience
for i, test_obs in enumerate(test_diff['Count'][200:400]):

    # build a model using the current buffers for endogenous and exogenous variables
    model = ARIMA(history_endog, exog=history_exog, order=(0, 0, 1)).fit(solver="bfgs", disp=0)

    # forecast the value for the test instance, supplying corresponding exogenous variables
    yhat = model.forecast(exog=test_exog.iloc[i])[0][0]

    # remember the forecasted value
    predictions.append(yhat)

    # update the buffers for the endogenous and exogenous variables
    history_endog.append(test_obs)
    history_exog.append(test_exog.iloc[i])

    print('predicted=%f, expected=%f' % (yhat, test_obs))
```

predicted=8.049571, expected=1.000000  
predicted=7.045170, expected=14.000000  
predicted=-1.095469, expected=2.000000  
predicted=2.527754, expected=-16.000000  
predicted=29.947982, expected=-2.000000  
predicted=34.744400, expected=27.000000  
predicted=12.712235, expected=-16.000000  
predicted=29.674745, expected=23.000000  
predicted=12.118860, expected=8.000000  
predicted=12.369147, expected=-29.000000  
predicted=44.930130, expected=-10.000000  
predicted=63.297150, expected=11.000000  
predicted=59.220923, expected=6.000000  
predicted=59.278726, expected=-1.000000  
predicted=67.233368, expected=-16.000000  
predicted=82.276874, expected=0.000000  
predicted=74.371741, expected=5.000000  
predicted=60.779669, expected=26.000000  
predicted=32.726741, expected=-10.000000  
predicted=39.015032, expected=-15.000000  
predicted=63.922744, expected=0.000000  
predicted=65.686789, expected=-2.000000  
predicted=64.652376, expected=-6.000000  
predicted=64.983977, expected=5.000000  
predicted=56.406325, expected=23.000000  
predicted=27.717255, expected=-12.000000  
predicted=44.362293, expected=-5.000000  
predicted=24.860091, expected=0.000000  
predicted=20.205686, expected=-2.000000  
predicted=14.317995, expected=8.000000  
predicted=6.451927, expected=-5.000000  
predicted=9.048683, expected=-10.000000  
predicted=14.851674, expected=4.000000  
predicted=13.449953, expected=6.000000  
predicted=9.409727, expected=5.000000  
predicted=8.971273, expected=-7.000000  
predicted=14.045981, expected=-6.000000  
predicted=11.582802, expected=-2.000000  
predicted=7.628140, expected=1.000000  
predicted=2.850881, expected=12.000000  
predicted=-3.178514, expected=1.000000  
predicted=-1.437983, expected=4.000000  
predicted=-1.673059, expected=-20.000000  
predicted=8.167300, expected=-1.000000  
predicted=8.840056, expected=3.000000  
predicted=5.188164, expected=30.000000  
predicted=-1.120441, expected=-24.000000  
predicted=13.837212, expected=21.000000  
predicted=-2.671490, expected=-18.000000  
predicted=10.856836, expected=19.000000  
predicted=8.675168, expected=-30.000000  
predicted=22.276658, expected=9.000000  
predicted=11.951569, expected=4.000000  
predicted=3.570007, expected=17.000000  
predicted=-1.132803, expected=-10.000000  
predicted=6.780990, expected=2.000000  
predicted=4.116245, expected=-7.000000  
predicted=8.282302, expected=1.000000  
predicted=5.588207, expected=0.000000  
predicted=8.211904, expected=-7.000000  
predicted=10.851715, expected=1.000000

predicted=5.326752, expected=-4.000000  
predicted=5.900269, expected=-7.000000  
predicted=6.011125, expected=17.000000  
predicted=-0.497053, expected=17.000000  
predicted=2.901083, expected=-28.000000  
predicted=9.744776, expected=-1.000000  
predicted=3.651410, expected=14.000000  
predicted=-1.896352, expected=-15.000000  
predicted=5.188745, expected=-4.000000  
predicted=3.627529, expected=35.000000  
predicted=-7.994481, expected=-20.000000  
predicted=4.572872, expected=2.000000  
predicted=1.629642, expected=-15.000000  
predicted=6.459738, expected=29.000000  
predicted=-2.723809, expected=-14.000000  
predicted=6.255069, expected=-5.000000  
predicted=6.934033, expected=-5.000000  
predicted=5.650863, expected=-7.000000  
predicted=7.090275, expected=27.000000  
predicted=-3.779565, expected=-5.000000  
predicted=0.932110, expected=-2.000000  
predicted=4.379474, expected=-9.000000  
predicted=5.611193, expected=-6.000000  
predicted=8.522218, expected=-1.000000  
predicted=3.742516, expected=9.000000  
predicted=-0.774481, expected=1.000000  
predicted=0.281647, expected=1.000000  
predicted=3.111240, expected=-10.000000  
predicted=4.262043, expected=-3.000000  
predicted=2.552074, expected=8.000000  
predicted=-1.341889, expected=8.000000  
predicted=-1.572691, expected=-2.000000  
predicted=0.646024, expected=-12.000000  
predicted=4.008323, expected=-4.000000  
predicted=2.324386, expected=0.000000  
predicted=1.818305, expected=5.000000  
predicted=1.416685, expected=6.000000  
predicted=0.936692, expected=-3.000000  
predicted=3.255455, expected=-1.000000  
predicted=3.831058, expected=0.000000  
predicted=1.977145, expected=3.000000  
predicted=-0.000152, expected=1.000000  
predicted=2.070849, expected=-8.000000  
predicted=5.953171, expected=24.000000  
predicted=-3.981675, expected=-16.000000  
predicted=4.027551, expected=0.000000  
predicted=1.187689, expected=-11.000000  
predicted=6.215612, expected=3.000000  
predicted=4.640704, expected=-2.000000  
predicted=2.817852, expected=24.000000  
predicted=-5.588319, expected=-25.000000  
predicted=5.838106, expected=3.000000  
predicted=5.144255, expected=11.000000  
predicted=1.902667, expected=-5.000000  
predicted=6.264333, expected=10.000000  
predicted=0.714582, expected=-2.000000  
predicted=1.387055, expected=-17.000000  
predicted=5.535654, expected=1.000000  
predicted=3.695698, expected=5.000000  
predicted=2.646542, expected=-6.000000  
predicted=7.221542, expected=3.000000

predicted=2.591847, expected=22.000000  
predicted=-0.671676, expected=-22.000000  
predicted=6.111219, expected=5.000000  
predicted=1.025851, expected=-4.000000  
predicted=3.229125, expected=6.000000  
predicted=1.777474, expected=-4.000000  
predicted=4.051500, expected=0.000000  
predicted=3.554556, expected=10.000000  
predicted=0.658965, expected=-14.000000  
predicted=6.376260, expected=17.000000  
predicted=-2.261599, expected=-5.000000  
predicted=2.913733, expected=-1.000000  
predicted=3.064173, expected=-7.000000  
predicted=3.528791, expected=-6.000000  
predicted=3.475846, expected=7.000000  
predicted=-0.851155, expected=13.000000  
predicted=-0.251717, expected=23.000000  
predicted=-5.800708, expected=-29.000000  
predicted=10.301352, expected=2.000000  
predicted=7.459642, expected=13.000000  
predicted=0.426673, expected=-3.000000  
predicted=2.571195, expected=-19.000000  
predicted=8.446308, expected=10.000000  
predicted=3.755651, expected=4.000000  
predicted=1.263788, expected=2.000000  
predicted=0.020201, expected=-2.000000  
predicted=2.526459, expected=-19.000000  
predicted=9.598170, expected=-1.000000  
predicted=6.374635, expected=29.000000  
predicted=-2.355041, expected=-10.000000  
predicted=4.330081, expected=-11.000000  
predicted=4.716641, expected=-8.000000  
predicted=7.361235, expected=3.000000  
predicted=4.330854, expected=4.000000  
predicted=4.868279, expected=3.000000  
predicted=1.001749, expected=2.000000  
predicted=1.225399, expected=-10.000000  
predicted=3.695615, expected=4.000000  
predicted=0.549966, expected=-4.000000  
predicted=1.579874, expected=7.000000  
predicted=0.047254, expected=-10.000000  
predicted=3.139583, expected=0.000000  
predicted=2.037813, expected=0.000000  
predicted=0.672286, expected=6.000000  
predicted=-1.013910, expected=18.000000  
predicted=-3.967804, expected=9.000000  
predicted=-0.713985, expected=-21.000000  
predicted=9.253132, expected=-5.000000  
predicted=4.331296, expected=7.000000  
predicted=-0.598177, expected=17.000000  
predicted=-4.295485, expected=-30.000000  
predicted=7.295338, expected=9.000000  
predicted=1.328814, expected=5.000000  
predicted=0.589366, expected=-6.000000  
predicted=4.640042, expected=-8.000000  
predicted=3.283968, expected=-1.000000  
predicted=2.125759, expected=15.000000  
predicted=1.955681, expected=2.000000  
predicted=1.269640, expected=11.000000  
predicted=-1.788685, expected=-14.000000  
predicted=3.532327, expected=-4.000000

```

predicted=3.224179, expected=-8.000000
predicted=4.355228, expected=2.000000
predicted=2.753483, expected=1.000000
predicted=0.908368, expected=22.000000
predicted=-5.008030, expected=-21.000000
predicted=4.510716, expected=-4.000000
predicted=4.513868, expected=0.000000
predicted=3.447257, expected=16.000000
predicted=-1.142146, expected=-1.000000
predicted=2.778545, expected=-1.000000
predicted=1.515870, expected=-13.000000
predicted=4.161721, expected=8.000000
predicted=2.474644, expected=-2.000000
predicted=3.901164, expected=15.000000
predicted=-1.231352, expected=-9.000000
predicted=2.963568, expected=0.000000
predicted=1.587609, expected=-11.000000

```


- Plot predictions on a graph to compare with the observed values

In [102]:

```
pd.DataFrame({"true": testdiff['Count'][200:400], "predicted": predictions}).plot(figsize=(16, 4))
```

Out[102]:

<matplotlib.axes.\_subplots.AxesSubplot at 0x7f9151988310>



- The predicted values from the ARIMAX model matches the peaks and troughs more closely than the previous ARIMA model created. We can see that the sudden spikes are also predicted by the model.

In [103]:

```
mse = mean_squared_error(testdiff['Count'][200:400], predictions)
np.sqrt(mse)
```

Out[103]:

18.10303545197831



In [104]:

```
mda = get_mda(testdiff['Count'][200:400], predictions)
mda
```

Out[104]:

0.6733668341708543

The numerical measures of accuracy, RMSE score and the MDA score

## Model Evaluation

### RMSE Results

In [105]:

```
rmse_table = pd.DataFrame({
    'method': ['Persistence baseline', 'ARIMA', 'ARIMAX'],
    'RMSE': [20.72135517709077, 12.520498573775068, 18.10303545197831],
})

rmse_table.set_index("method", inplace=True)

# add columns with percent changes on the baselines
rmse_table['% Change on Pers. baseline'] = 100 * rmse_table['RMSE'] / 20.72135517709077 - 100

rmse_table
```

Out[105]:

	RMSE	% Change on Pers. baseline
method		
Persistence baseline	20.721355	0.000000
ARIMA	12.520499	-39.576835
ARIMAX	18.103035	-12.635852

The table above shows the root mean square error for the previously calculated baseline, ARIMA model and ARIMAX model. It shows that the ARIMA model produced the best rates: 39% reduction on the persistence baseline.

### MDA Results

In [106]:

```
mda_table = pd.DataFrame({
    'method': ['Persistence baseline', 'ARIMA', 'ARIMAX'],
    'MDA': [0.318475073313783, 0.7085427135678392, 0.6733668341708543 ]
})

mda_table.set_index("method", inplace=True)

# add columns with percent changes on the baselines
mda_table['% Change on Pers. baseline, MDA'] = 100 * mda_table['MDA'] / 0.318475073313783 - 100
mda_table
```

Out[106]:

	MDA	% Change on Pers. baseline, MDA
method		
Persistence baseline	0.318475	0.000000
ARIMA	0.708543	122.479802
ARIMAX	0.673367	111.434706

In terms of MDA, ARIMA model also produced the better results with the MDA increase of 122% on the persistence baseline

## Conclusion

Time series models; ARIMA and ARIMAX were used to forecast the number of accidents to help a road assistance company effectively allocate their resources to meet demand. From the EDA, it was shown that the number of accidents were higher between 7:30am - 8:30 and 3:30pm-6:30pm from Monday to Friday. It was said that this could be explained by the rush hour/peak time travel.

The ARIMAX model built included the exogenous variable `Urban` determining the effect the type of area had on the number of accidents that occurred.

Both models were used to predict instances of the test data where the ARIMAX model was able to match the peaks and troughs in the data much closer than the ARIMA model.

It can be said that as the predicted values matched the observed values, it can be said that the road assistance would be recommended to:

- Increase the number of resources allocated during the hours of rush hour/peak times
- Have a certain number of resources on standby for the remaining quieter hours of the day

**Possible further improvements**

- Though understanding that it is difficult to predict human error, would try more methods where time series models and other traditional ML models could be used simultaneously for better performance
- Using new data from most recent years, using it as new test set
- Take into consideration the effect of other variables on the number of accidents
- Determine whether holiday periods such as Christmas have an effect on the number of accidents

In [ ]: